

Intro to Quest

Northwestern IT Research Computing and
Data Services

Andrew Block, Wenceslao Hernandez Velasquez, Julia Giannini

June 26, 2025

Overview

1. Introduction
2. Navigating Quest
3. Ways to interact with Quest
4. Data Transfers
5. How to submit jobs
6. Quest OnDemand



Technical Support

quest-help@northwestern.edu



Requesting Help

- Request a consult or open a ticket!
 - quest-help@northwestern.edu
- Office hours:
 - Every Monday, 3-4 PM
 - Rooms 2202-2205, Mudd Library (2nd floor)
 - Across the outside hallway from the GIS Lab

What is High Performance Computing (HPC)?

High Performance Computing

HPC uses supercomputers and computer clusters to solve advanced computation problems. – Wikipedia

HPC encompasses solutions that are able to process data and execute calculations at a rate that far exceeds other computers. This aggregate computing power enables [scientists] to solve large problems that would otherwise be unapproachable. – HPE

HPC is technology that uses clusters of powerful processors, working in parallel, to process massive multi-dimensional datasets (big data) and solve complex problems at extremely high speeds. - IBM


Quest!



HPC at Northwestern

- Host big databases and facilitate computing for astrophysics, biology, etc.
- Enable Genomics Research and Pipelines with GCC
- Be able to use the power of multiple computers
- Quest:
 - 5,500 active users, 1231 compute nodes, 80,390 CPU cores, 304 GPU Cards, 750 research software, 78 knowledge base articles, 28 training recordings

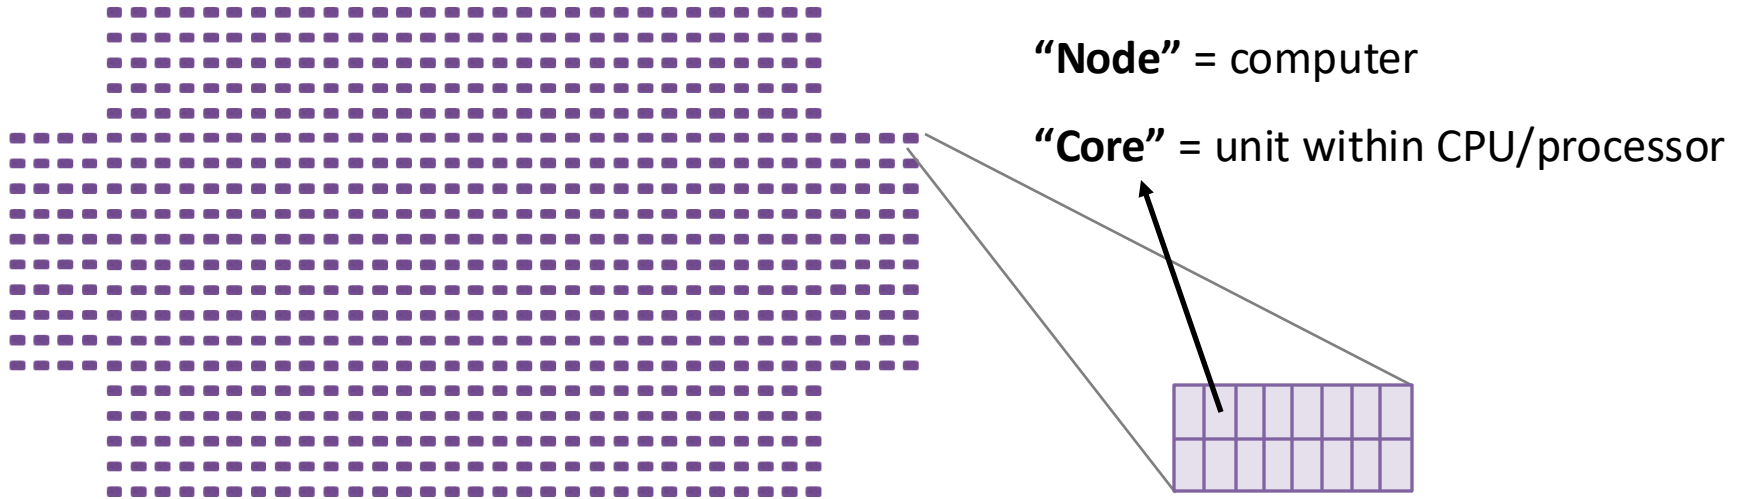
Quest consists of ~1200 nodes



→ **“Node”** = computer




Each node consists of 52-128 cores





Each node can run separate jobs on each core

Python	Fortran	R	MATLAB	MATLAB	Vasp	R	Python
GATK	GATK	GATK	Python	R	Perl	MATLAB	Python


 = multi-core job

A node can run a job with 52-128 cores at a time




 = one multi-core job

You can run thousands of jobs at a time



 = a single job

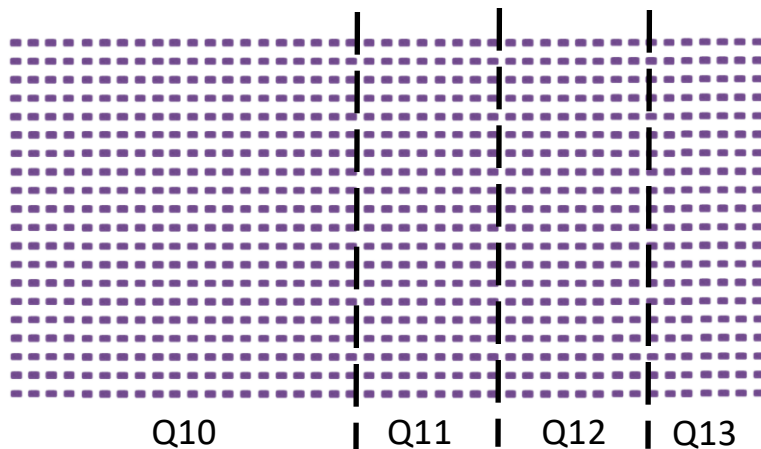
Jobs can even run across multiple nodes



 = a single job

Technical specifications of nodes vary

There are several [generations of nodes](#)!




Generation	# nodes	# cores per node	Total Schedulable memory (RAM)	Memory per core
Q10	555	52	166 GB	~3.2 GB
Q11	209	64	221 GB	~3.5 GB
Q12	214	64	221 GB	~3.5 GB
Q13	140	128	473 GB	~3.7GB
<hr/>				
Your laptop (probably)	1	10-25	16-128GB	~2-8GB

Computing Resources



Each *node* has *schedulable resources*:

- 40-128 *CPU cores*
- 192-512 GB of *CPU RAM* (total)
- 0-4 *GPU cards*
- 40-80 GB of *GPU RAM* (per card)



Specialty (general access) compute nodes on Quest

- High memory node available through General Access
 - 1.5 TB of RAM
 - 52 CPU cores
- GPU nodes available through General Access --> parallelization & speed
 - 58 nodes with 2 or 4 GPU cards each



Quest technical details

- Cluster network interconnect: Infiniband
 - High speed communication between nodes, storage, etc.
- Storage: General Parallel File System (GPFS) ≈ 8PB (1PB = 10^{15} B)
- Operating system: Red Hat Enterprise Linux (RHEL) 8.9
- Scheduler: SLURM (Simple Linux Utility or Resource Management)




Overview

1. Introduction
2. Navigating Quest
3. Ways to interact with Quest
4. Data Transfers
5. How to submit jobs
6. Quest OnDemand

Allocations on Quest

- An allocation is a project group that allows you to do research and submit jobs on Quest
 - Provides:
 - Storage (personal and shared) space
 - Compute resources and the ability to submit jobs
- Types of allocations
 - Research I
 - Research II
 - Buy-in
 - Classroom



How do I access files and folders on Quest?

- **File system is shared across all computers on Quest**

→ You can access your files & folders from anywhere

- **You will be working with 3 main folders:**

- `/home/<netid>` ← your personal home directory
- `/projects/<alloc-id>` ← your project/allocation directories
- `/hpc/software` ← where software modules live

-
- `/scratch /<netid>` ← 5TB of temporary (30 day) storage

/home & /projects directories on Quest

/home/<netid>

/projects/<allocation-id>

Who has access

- Just you
- Everyone in the allocation

Storage

- 80GB storage space
- Much more space

Contains

- Software, codes etc.
- Data, outputs

Check for space

- “homedu”
- “checkproject <alloc-id>”

Backups

- Backed up every 24 hrs
- NOT BACKED UP

Who will have access to my files?



Files created in
/projects



- Read/write permissions to all allocation members by default

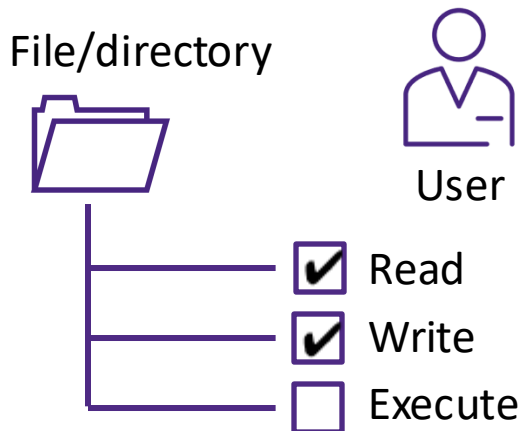


Files created in
/home



- Read/write permissions to only you by default

What are “permissions”?




Files you created:

- You will always have full read/write permissions.

Files you did not create:

- You may not have read/write permissions.

Group-level permissions




Group permissions

- Permission are also defined at the “user group” level.
- To see which user groups you are a part of, run “groups”

```
[netid@quser21 ~]$ groups  
netid      p30XXX    b10XX
```


Permission strings

```
[abc1234@quser21 p12345]$ ls -l
total 2
drwxrwxr-x 1 abc1234 p12345 4096 Apr 15 11:00 codes
drwxrwxr-x 1 abc1234 p12345 4096 Apr 15 10:00 data
-rw-rw-r-- 1 abc1234 p12345 259 Apr 15 09:00 test.txt
```



A bit more on permissions:

<https://kb.northwestern.edu/70712>



Overview

1. Introduction
2. Navigating Quest
3. Ways to interact with Quest
4. Data Transfers
5. How to submit jobs
6. Quest OnDemand

Ways to interact with Quest

- Login nodes
 - For testing/development purposes can run things here
 - Limited to 4 CPUs and 4GB RAM
 - Submit **batch** or **interactive** jobs (terminal)
- Quest Analytics (web interface for interactive applications)
- Quest OnDemand (mainly **interactive** jobs)

```
#####  
##  
## (QUEST) ##  
##  
##  
##  
##  
##  
#####  
  
Welcome to Northwestern's high-performance compute cluster! Information on Quest can be found at https://www.it.northwestern.edu/departments/it-services-support/research/computing/quest/index.html. For support or to purchase nodes or storage for dedicated usage, please contact Research Computing Services at quest-help@northwestern.edu. Quest's Storage and Data Policy is available at https://www.it.northwestern.edu/departments/it-services-support/research/computing/quest/storage-data-policy.html.  
  
Quest is not approved for contractually or legally restricted data. Ensure sensitive data is anonymized or de-identified before being moved to Quest. Email quest-help@northwestern.edu if you need a computing environment for the use of restricted data.
```

Logging into Quest - Terminal

- Make sure you're connected to [NU VPN](#)
- Log in options
 - [Terminal – SSH](#)
`ssh <net_id>@login.quest.northwestern.edu`
 - Windows: use a terminal with Linux commands

Ways to interact with Quest

- [Analytics nodes](#) (browser)
 - Run interactive Jupyter, R, or SAS sessions
 - Note usage guidelines on our website
 - Processes will automatically be killed if they overutilize CPUs or memory

[Logout](#)[Control Panel](#)[Files](#)[Running](#)[Clusters](#)

Select items to perform actions on them.

[Upload](#)[New ▾](#)☐ 0 ▾ /[Name ▾](#)[Last Modified](#)[File size](#)☐ [bin](#)

21 days ago

☐ [Downloads](#)

8 days ago

Logging into Quest – Quest OnDemand

- [Quest OnDemand](#) (browser)
- Interactive jobs and GUIs
 - R, Jupyter, VSCode, Matlab, Mathematica, Stata, and more!

Quest OnDemand Apps ▾ Files ▾ Jobs ▾ Clusters ▾ Interactive Apps ▾   ▾  ▾ Request a Consultation ▾

Northwestern | INFORMATION TECHNOLOGY
RESEARCH COMPUTING AND DATA SERVICES

Northwestern IT Research Computing and Data Services presents Quest OnDemand, a single access point for using Quest, Northwestern's High-Performance Compute Cluster.

Overview

1. Introduction
2. Navigating Quest
3. Ways to interact with Quest
4. Data Transfers
5. How to submit jobs
6. Quest OnDemand

Data Transfers

- Ways to transfer your data:
 - Globus
 - Quest OnDemand uploads
 - Command line utilities (`sftp`, `scp`, `rsync`, `wget`, `git`)


Globus

- Globus completes data transfers between data storage locations and compute systems quickly and securely
 - Includes a web interface and a command line interface (CLI)
- Automated transfers
- Robust to connectivity problems
 - If you lose connection, Globus will continue where you left off
- [Knowledge Base article](#)

Globus Endpoints

- Endpoint: A location where authorized users can transfer data to and from such as:

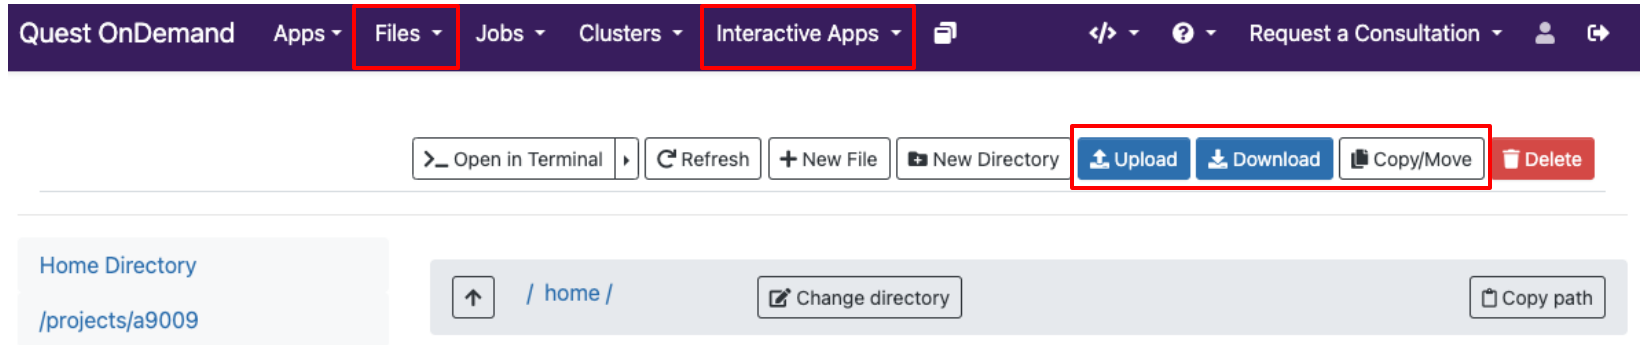
- Quest
- Your laptop
- Your collaborator's laptops
- Other university's HPCs



Globus demo

Quest OnDemand

- Web-based interface to run interactive applications on Quest compute nodes
- File explorer:
 - View, edit, and upload/download files
- More on Quest OnDemand later!



File transfer through the command line

- **Command line utilities:** `sftp`, `scp`, `rsync`, `wget`, `git`
- **SFTP – Secure File Transfer Protocol**
 - Protects data from unauthorized access
 - `$ sftp <netid>@login.quest.northwestern.edu` – **establish connection.**
 - **Upload files:** `$ put <source_file> <destination_file>`
 - **Downloading files:** `$ get <source_file> <destination_file>`

Overview

1. Introduction
2. Navigating Quest
3. Ways to interact with Quest
4. Data Transfers
5. How to submit jobs
6. Quest OnDemand

Quiz!

Break

Overview

1. Introduction
2. Navigating Quest
3. Ways to interact with Quest
4. Data Transfers
5. How to submit jobs
6. Quest OnDemand

Types of jobs: batch vs. interactive

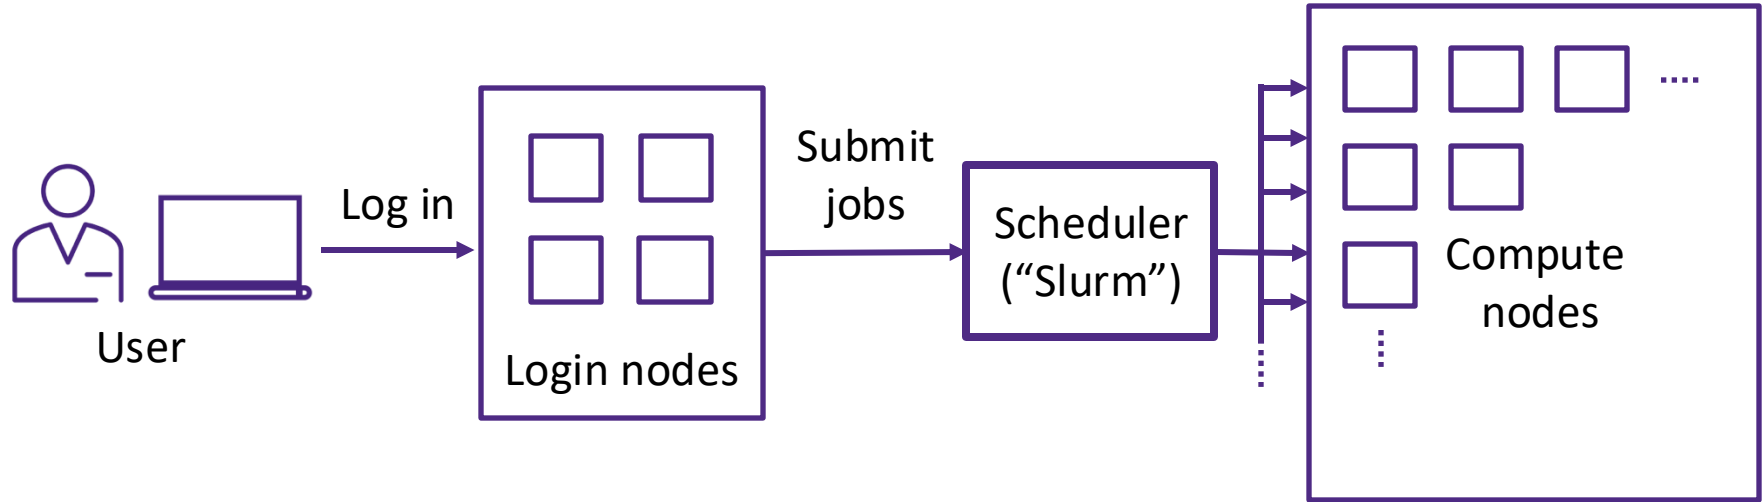
- **Batch job**

- Submit your job as a pre-written bash script
- Benefit – submit & forget about it, high throughput

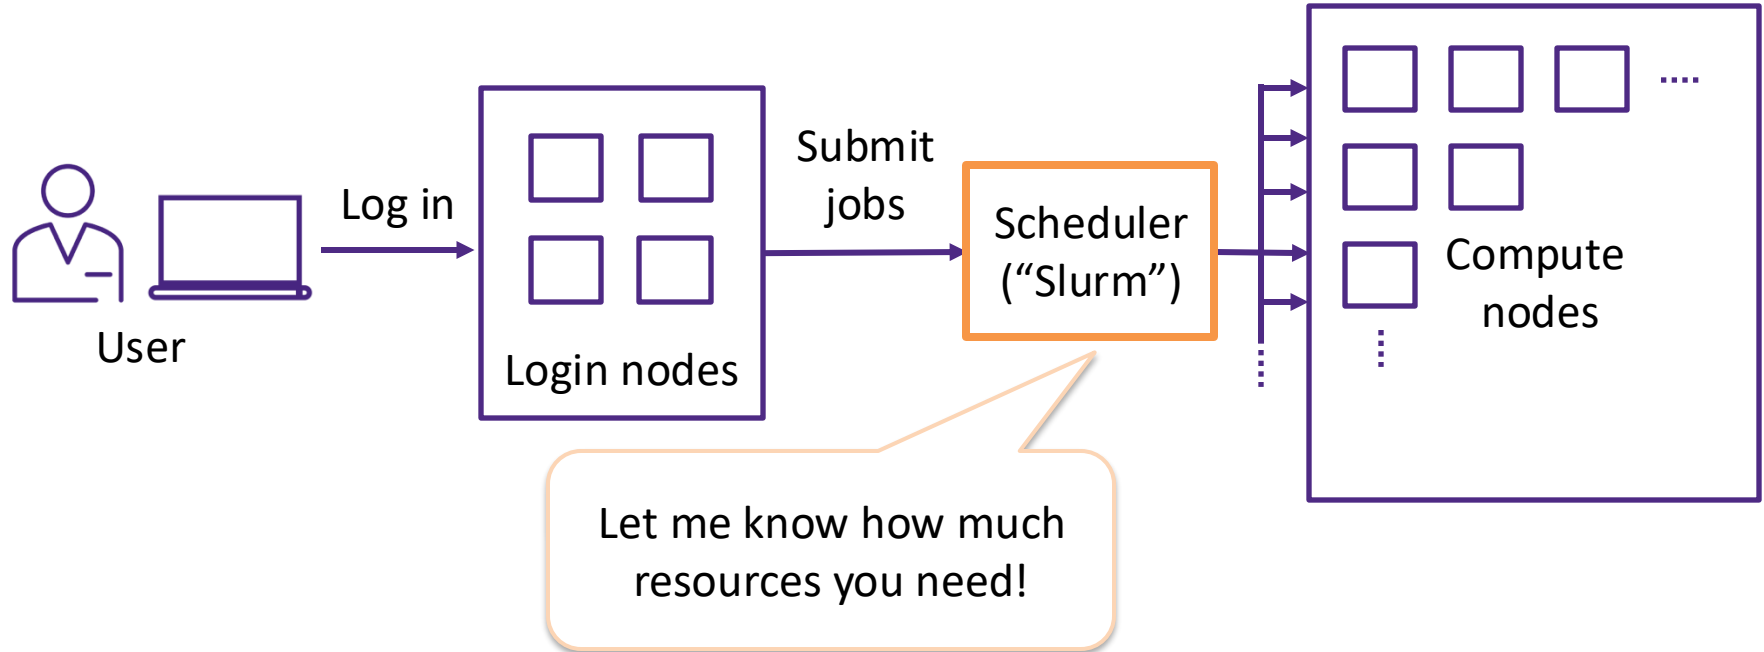
- **Interactive job**

- Run interactive session on the *compute nodes*
- Benefit – exploratory work, troubleshooting etc.

Log into the login nodes & submit jobs to the scheduler



Log into the login nodes & submit jobs to the scheduler



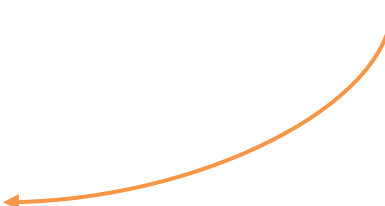
What does a job script look like?

```
#!/bin/bash
#SBATCH --account=p12345
#SBATCH --partition=short
#SBATCH --nodes=1
#SBATCH --ntasks-per-node=1
#SBATCH --time=00:10:00
#SBATCH --mem=3G
#SBATCH --job-name=sample_job
#SBATCH --output=outlog

## load modules (python)
module purge all
module load python

## run my program
python hello_world.py
```

Your script might look like this.
Let's break down the components!



What does a job script look like?

```
#!/bin/bash
#SBATCH --account=p12345
#SBATCH --partition=short
#SBATCH --nodes=1
#SBATCH --ntasks-per-node=1
#SBATCH --time=00:10:00
#SBATCH --mem=3G
#SBATCH --job-name=sample_job
#SBATCH --output=outlog
```

```
## load modules (python)
module purge all
module load python
```

```
## run my program
python hello_world.py
```

Tell Slurm what you want
with these “headers”

Load any modules you need

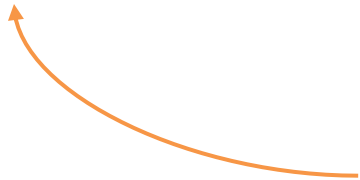
Run your cool program!
(can be multiple commands)

Side note: there are short-forms for headers

```
#!/bin/bash
#SBATCH --account=p12345
#SBATCH --partition=short
#SBATCH --nodes=1
#SBATCH --ntasks-per-node=1
#SBATCH --time=00:10:00
#SBATCH --mem=3G
#SBATCH --job-name=sample_job
#SBATCH --output=outlog
```

=

```
#!/bin/bash
#SBATCH -A p12345
#SBATCH -p short
#SBATCH -N 1
#SBATCH --ntasks-per-node=1
#SBATCH -t 00:10:00
#SBATCH --mem=3G
#SBATCH -J sample_job
#SBATCH --output=outlog
```



We will use the long-form in
this workshop for clarity!

Indicate allocation in “--account” or “-A”

```
#!/bin/bash
#SBATCH --account=p12345
#SBATCH --partition=short
#SBATCH --nodes=1
#SBATCH --ntasks-per-node=1
#SBATCH --time=00:10:00
#SBATCH --mem=3G
#SBATCH --job-name=sample_job
#SBATCH --output=outlog

## load modules (python)
module purge all
module load python

## run my program
python hello_world.py
```

Which allocation should I use?

→ Try running “groups <netid>”

```
[abc1234@quser21 ~]$ groups abc1234
abc1234 p12345 b1000 e10001
```

Indicate partition in “--partition” or “-p”

```
#!/bin/bash
#SBATCH --account=p12345
#SBATCH --partition=short
#SBATCH --nodes=1
#SBATCH --ntasks-per-node=1
#SBATCH --time=00:10:00
#SBATCH --mem=3G
#SBATCH --job-name=sample_job
#SBATCH --output=outlog

## load modules (python)
module purge all
module load python

## run my program
python hello_world.py
```

How do I know which partition to select?

If your allocation starts with “p” or “e”:

- “short” (4h) “normal” (48h) “long” (168 h)
- “gengpu” for GPU
- “genhimem” for high-memory

If your allocation starts with “b”

- generally, use allocation name (some special cases)
e.g. “b1234”

Indicate number of nodes in “--nodes” or “-N”

```
#!/bin/bash
#SBATCH --account=p12345
#SBATCH --partition=short
#SBATCH --nodes=1
#SBATCH --ntasks-per-node=1
#SBATCH --time=00:10:00
#SBATCH --mem=3G
#SBATCH --job-name=sample_job
#SBATCH --output=outlog

## load modules (python)
module purge all
module load python

## run my program
python hello_world.py
```

How many nodes should I select?

Unless your program explicitly states that it can run across nodes, default to nodes=1.

If your job doesn't rely on software like **OpenMPI, MPICH, Intel-MPI**, then most likely nodes=1.

Indicate number of cores in “--ntasks-per-node”

```
#!/bin/bash
#SBATCH --account=p12345
#SBATCH --partition=short
#SBATCH --nodes=1
#SBATCH --ntasks-per-node=1
#SBATCH --time=00:10:00
#SBATCH --mem=3G
#SBATCH --job-name=sample_job
#SBATCH --output=outlog

## load modules (python)
module purge all
module load python

## run my program
python hello_world.py
```

How many cores should I select?

Unless your program is specifically designed to run on multiple cores, select “1”.

Examples of multi-core:

Python’s scikit-learn, R’s DESEQ2 etc.

NOTE: you still need to explicitly tell the software!

Indicate time in “--time” or “-t”

```
#!/bin/bash
#SBATCH --account=p12345
#SBATCH --partition=short
#SBATCH --nodes=1
#SBATCH --ntasks-per-node=1
#SBATCH --time=00:10:00
#SBATCH --mem=3G
#SBATCH --job-name=sample_job
#SBATCH --output=outlog

## load modules (python)
module purge all
module load python

## run my program
python hello_world.py
```

How much time should I request?

- Start with a conservative estimate.
- Once you have a better idea of how much time your job will take, adjust accordingly.

Indicate memory in “--mem” or “--mem-per-cpu”

```
#!/bin/bash
#SBATCH --account=p12345
#SBATCH --partition=short
#SBATCH --nodes=1
#SBATCH --ntasks-per-node=1
#SBATCH --time=00:10:00
#SBATCH --mem=3G
#SBATCH --job-name=sample_job
#SBATCH --output=outlog

## load modules (python)
module purge all
module load python

## run my program
python hello_world.py
```

How much memory should I request?

Be careful as this affects your job's pending time by a lot!

- If you request a lot of memory, your job might be restricted to certain nodes.
- You can take the same approach and start with a conservative estimate, but make sure to adjust later.

Why not request maximum time and memory?

- **On Quest, job priority is determined using the user's FairShare Score**
 - FairShare Score makes sure that everyone gets their fair share
 - Score determined based on the amount of resources you request, not the amount that you end up using
 - Requesting a ton of resources will lower your job priority and increase pending time
 - HPC clusters consume a lot of energy
 - Strive for computational efficiency 🌱

How to check resource utilization

Useful commands:

- `checkjob <job_id>` - detailed summary
- `seff <job_id>` - efficiency summary

Can you spot any problematic trends that should be adjusted?

```
[quest_demo@quser23]$ seff 2261088
-----
JOB EFFICIENCY
-----
Job ID: 2261088
Cluster: quest
User/Group: netid/netid
State: COMPLETED (exit code 0)
Cores: 4
CPU Utilized: 00:32:42
CPU Efficiency: 24.84% of 00:32:42 core-walltime
Job Wall-clock time: 00:32:42
Memory Utilized: 366.84 MB
Memory Efficiency: 3.66% of 10.00 GB
```


How to check resource utilization

Useful commands:

- `checkjob <job_id>` - detailed summary
- `seff <job_id>` - efficiency summary

Can you spot any problematic trends that should be adjusted?

```
[quest_demo@quser43]$ seff 2261088
```

```
-----  
JOB EFFICIENCY  
-----
```

```
Job ID: 2261088
```

```
Cluster: quest
```

```
User/Group: netid/netid
```

```
State: COMPLETED (exit code 0)
```

```
Cores: 4
```

```
CPU Utilized: 00:32:42
```

```
CPU Efficiency: 24.84% of 00:32:42 core-walltime
```

```
Job Wall-clock time: 00:32:42
```

```
Memory Utilized: 366.84 MB
```

```
Memory Efficiency: 3.66% of 10.00 GB
```

Set job name for ease of identification

```
#!/bin/bash
#SBATCH --account=p12345
#SBATCH --partition=short
#SBATCH --nodes=1
#SBATCH --ntasks-per-node=1
#SBATCH --time=00:10:00
#SBATCH --mem=3G
#SBATCH --job-name=sample_job
#SBATCH --output=outlog

## load modules (python)
module purge all
module load python

## run my program
python hello_world.py
```

How to decide job name?

This is for your own convenience – Slurm will not care what you name your job. Something short and descriptive can be useful later.

Set output file path

```
#!/bin/bash
#SBATCH --account=p12345
#SBATCH --partition=short
#SBATCH --nodes=1
#SBATCH --ntasks-per-node=1
#SBATCH --time=00:10:00
#SBATCH --mem=3G
#SBATCH --job-name=sample_job
#SBATCH --output=outlog
```

```
## load modules (python)
module purge all
module load python
```

```
## run my program
python hello_world.py
```

How to decide output file path?

- Slurm expects you to specify a **full path** with a file name, not just a directory.
- Direct your output to **/projects**, not **/home**! Some programs write a lot of output and can overflow your **/home**

Load software into the environment

```
#!/bin/bash
#SBATCH --account=p12345
#SBATCH --partition=short
#SBATCH --nodes=1
#SBATCH --ntasks-per-node=1
#SBATCH --time=00:10:00
#SBATCH --mem=3G
#SBATCH --job-name=sample_job
#SBATCH --output=outlog

## load modules (python)
module purge all
module load python

## run my program
python hello_world.py
```



Load any modules you need

Accessing/managing software on Quest

- [Module system](#)
 - Search for a module: `module spider <name_of_software>`
- **Virtual environments**
 - [Anaconda virtual environments on Quest](#)
- Containers
- Compile/install code from source
- [Request a software installation](#)

```
[abc1234@quser31 ~]$ module spider python
```

```
python:
```

```
-----  
Versions:
```


```
python/2.7.5  
python/2.7.18-gcc  
python/3.6.10  
python/3.8.12-gcc  
python/3.8.12-intel  
python/3.10.1  
...
```

Execute commands/scripts

```
#!/bin/bash
#SBATCH --account=p12345
#SBATCH --partition=short
#SBATCH --nodes=1
#SBATCH --ntasks-per-node=1
#SBATCH --time=00:10:00
#SBATCH --mem=3G
#SBATCH --job-name=sample_job
#SBATCH --output=outlog

## load modules (python)
module purge all
module load python

## run my program
python hello_world.py
```



Run your cool program!
(can be multiple commands)

Submit a batch job with sbatch

- `sbatch <name_of_submit_script>`

```
[abc1234@quser31 ~]$ sbatch submit_script.sh
```

```
Submitted batch job 7342820
```

Types of jobs: batch vs. interactive

- **Batch job**

- Submit your job as a pre-written bash script
- Benefit – submit & forget about it

- **Interactive job**

- Run interactive session on the compute nodes
- Benefit – exploratory work, troubleshooting etc.

Interactive job resources are requested similarly

```
[netid@quser23]$ srun --account=p12345 --time=1:00:00 -n 1 -p short --mem=1G --pty bash -l
-----
srun job start: Wed Oct  5 14:11:43 CDT 2022
Job ID: XXXXXXXX
Username: netid
Queue: short
Account: p12345
-----
The following variables are not guaranteed to be the same in prologue and the job run
script
-----
PATH (in prologue) :
/usr/local/bin:/usr/bin:/usr/local/sbin:/usr/sbin:/usr/lpp/mmfs/bin:/opt/ibutils/bin:/opt/z
eek/bin:/home/netid/.local/bin:/home/netid/bin
WORKDIR is: /home/netid
-----
[netid@qnode8075]$ # code away!
```

“salloc” similar to “srun” but requires SSH

```
[quest_demo@quser23]$ salloc --x11 --account=p12345 --partition=short --nodes=1 --tasks-per-  
node=1 --time=00:20:00  
salloc: Pending job allocation 8575069  
salloc: job 8575069 queued and waiting for resources  
salloc: job 8575069 has been allocated resources  
salloc: Granted job allocation 8575069  
salloc: Waiting for resource configuration  
salloc: Nodes qnode6702 are ready for job  
[quest_demo@quser23]$ ssh -X qnode6702  
Last login: Mon Dec 13 12:46:35 2021 from quser23  
[quest_demo@qnode6702 ~]$ # Code away!
```

Useful commands 1/2: see pending and running jobs

```
[netid@quser23]$ squeue -u netid # inspect pending or running jobs
```

JOBID	PARTITION	NAME	USER	ST	TIME	NODES	NODELIST (REASON)
2632440	p12345	sample_j	tempuser	R	0:02	1	qnode5057
2632441	p12345	sample_j	tempuser	R	0:02	1	qnode5057

Useful commands 2/2:

see all jobs within a time window


```
[netid@quser23]$ sacct -X -u netid --starttime=9/24/22 --endtime=9/28/22 # past jobs
```

JobID	JobName	Partition	Account	AllocCPUS	State	ExitCode
1835975	testjob	short	p12345	1	TIMEOUT	0:0
2261088	testjob	normal	p12345	1	COMPLETED	0:0
2261164	testjob	normal	p12345	1	CANCELLED+	0:0

Documentation and example jobs

Everything You Need to Know About Using Slurm on Quest

<https://github.com/nuitrcs/examplejobs>



Exercise 1

Exercise

1. Using the right git command, pull down the repository from https://github.com/nuitrcs/REU_Activities to your computer
2. Use your text editor of choice to open `start_example_submit.sh`
3. Finish the submission script!



Work with a friend if you like!

Exercise - Answer

```
#!/bin/bash
#SBATCH --account=e32894  ## YOUR ACCOUNT pXXXX or bXXXX
#SBATCH --partition=short  ### PARTITION (buyin, short, normal, etc)
#SBATCH --nodes=1  ## how many computers do you need
#SBATCH --ntasks-per-node=1  ## how many cpus or processors do you need on each computer
#SBATCH --time=00:10:00  ## how long does this need to run (remember different partitions have
restrictions on this param)
#SBATCH --mem-per-cpu=1G  ## how much RAM do you need per CPU (this effects your FairShare score so
be careful to not ask for more than you need))
#SBATCH --job-name=sample_job  ## When you run squeue -u NETID this is how you can identify the job
#SBATCH --output=outlog  ## standard out and standard error goes to this file
#SBATCH --mail-type=ALL  ## you can receive e-mail alerts from SLURM when your job begins and when
your job finishes (completed, failed, etc)
#SBATCH --mail-user=email@u.northwestern.edu  ## your email

module purge all
module load python/3.12.10

python --version
python slurm_test.py
```


Overview

1. Introduction
2. Navigating Quest
3. Ways to interact with Quest
4. How to submit jobs
5. Quest OnDemand

Quest OnDemand Demo

Exercise 2

Thank You!

Questions?

quest-help@northwestern.edu



Request a
Consultation

<https://www.it.northwestern.edu/departments/it-services-support/research/computing/quest/>